

Model `rgeneric`

This is a class of generic models allows the user to define latent model-component in R, for cases where the requested model is not yet implemented in INLA, and do the Bayesian inference using INLA. It will run slower as the model properties has to be evaluated in R within a C-program.

Defining a latent model in R

The use of this feature, is in short the following. First we pass our definition of the model `rmodel`, to define a `inla-rgeneric` object,

```
model = inla.rgeneric.define(rmodel, debug, R.init, <args>)
```

Here, `rmodel` is model definition encoded as an R-function, `debug` is a logical parameter if debug information should be printed, `R.init` is a R-file which is `sourced` or `loaded` before the model is defined, and `<args>` are further parameters to the model (like dimension, prior-settings, etc). Then the model can used as

```
y ~ ... + f(i, model=model)
```

Example: the AR1 model

The `rmodel` needs to follow some rules to provide the required features. As an example, we will show how to implement the AR1-model, see `inla.doc("ar1")`. This model is defined as¹

$$x_1 \sim \mathcal{N}(0, \tau)$$

and

$$x_t \mid x_1, \dots, x_{t-1} \sim \mathcal{N}(\rho x_{t-1}, \tau_I), \quad t = 2, \dots, n.$$

The scale-parameter is the marginal precision τ , but the joint density is more naturally expressed using the innovation precision $\tau_I = \tau/(1 - \rho^2)$. The joint density of x is Gaussian

$$\pi(x \mid \rho, \tau) = \left(\frac{1}{\sqrt{2\pi}} \right)^n \tau_I^{n/2} (1 - \rho^2)^{1/2} \exp \left(-\frac{\tau_I}{2} x^T R x \right)$$

where the precision-matrix is

$$Q = \tau_I R = \tau_I \begin{bmatrix} 1 & -\rho & & & & \\ -\rho & 1 + \rho^2 & -\rho & & & \\ & -\rho & 1 + \rho^2 & -\rho & & \\ & & \ddots & \ddots & \ddots & \\ & & & -\rho & 1 + \rho^2 & -\rho \\ & & & & -\rho & 1 \end{bmatrix}$$

There are two (hyper-)parameters for this model, it is the marginal precision τ and the lag-one correlation ρ . We reparameterise these as

$$\tau = \exp(\theta_1)$$

and

$$\rho = 2 \frac{\exp(\theta_2)}{1 + \exp(\theta_2)} - 1$$

¹The second argument in $\mathcal{N}(,)$ is the precision not the covariance.

It is required that the parameters $\theta = (\theta_1, \theta_2)$ have support on \Re and the priors for τ and ρ are given as the corresponding priors for θ_1 and θ_2 . **IMPORTANT: A good re-parameterisation is required for INLA to work well. A good parameterisation makes, ideally, the “Fisher information matrix” of θ constant (wrt to θ). It is sufficient to check this in a frequentistic setting with data directly from the AR(1) model, in this case.** Note that INLA only provide the marginal posteriors for θ , but you can use `inla.tmarginal` to convert it to the appropriate marginals for ρ and τ .

We assign a (Gamma) $\Gamma(\cdot; a, b)$ prior (with mean a/b and variance a/b^2) for τ and a Gaussian prior $\mathcal{N}(\mu, \kappa)$ for θ_2 , so the joint prior for θ becomes

$$\pi(\theta) = \Gamma(\exp(\theta_1); a, b) \exp(\theta_1) \times \mathcal{N}(\theta_2; \mu, \kappa).$$

We will use $a = b = 1$, $\mu = 0$ and $\kappa = 1$.

In order to define the AR1-model, we need to make functions that returns

- the graph,
- the precision matrix $Q(\theta)$,
- the zero mean,
- the initial values of θ ,
- the log-normalising constant, and
- the log-prior

which except for the graph, depends on the current value of θ . We need to wrap this into a common function, which process the request from the C-program. The list of commands and its names

```
cmd = c("Q", "graph", "mu", "initial", "log.norm.const",
        "log.prior", "quit"),
```

are fixed. The skeleton-function for defining a model is

```
‘inla.rgeneric.ar1.model’ = function(
  cmd = c("graph", "Q", "mu", "initial", "log.norm.const",
          "log.prior", "quit"),
  theta = NULL,
  args = NULL)
{
  graph = function(n, theta){ <to be completed> }
  Q = function(n, theta) { <to be completed> }
  mu = function(n, theta) { <to be completed> }
  log.norm.const = function(n, theta) { <to be completed> }
  log.prior = function(n, theta) { <to be completed> }
  initial = function(n, theta) { <to be completed> }
  quit = function(n, theta) { <to be completed> }

  cmd = match.arg(cmd)
  val = do.call(cmd, args = list(n = as.integer(args$n), theta = theta))
  return (val)
}
```

The input parameters are

cmd What to return

theta The values of the θ -parameters

args User-defined argument(s) from the `inla.rgeneric.define()`-call; all arguments in `<args>` as a list. In this example we use `n` as the user-defined argument and pass it to every function. We then need to pass `n` when we define the model

```
model = inla.rgeneric.define(inla.rgeneric.ar1.model, n = 100)
```

This is just an example, and you may chose to do this differently.

Our next task, is the “fill in the blanks” in this function.

Function graph

This function must return a `sparseMatrix`, with the non-zero elements of the precision matrix. Only the lower-triangular part of the matrix is used.

```
graph = function(n, theta)
{
  ## return the graph of the model. the values of Q is only interpreted as zero or
  ## non-zero. return a sparse.matrix
  if (FALSE) {
    ## slow and easy: dense-matrices
    G = toeplitz(c(1, 1, rep(0, n-2L)))
    G = inla.as.sparse(G)
  } else {
    ## faster. we only need to define the lower-triangular of G
    i = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      1L:(n-1L))
    j = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      2L:n)
    x = 1 ## meaning that all are 1
    G = sparseMatrix(i=i, j=j, x=x, giveCsparse = FALSE)
  }
  return (G)
}
```

Function Q

This function must return the precision matrix $Q(\theta)$, and must be a `sparseMatrix`. Only the lower-triangular part of the matrix is used. We will make use of the helper function

```
interpret.theta = function(n, theta)
{
  ## internal helper-function to map the parameters from the internal-scale to the
  ## user-scale
  return (list(prec = exp(theta[1L]),
    rho = 2*exp(theta[2L])/(1+exp(theta[2L])) - 1.0))
}
```

to convert from θ_1 to τ , and from θ_2 to ρ . The `Q`-function can then be implemented as follows.

```
Q = function(n, theta)
{
  ## returns the precision matrix for given parameters
```

```

param = interpret.theta(n, theta)
if (FALSE) {
  ## slow and easy: dense-matrices
  Q = param$prec/(1-param$rho^2) * toeplitz(c(1+param$rho^2, -param$rho, rep(0, n-2L)))
  Q[1, 1] = Q[n, n] = param$prec/(1-param$rho^2)
  Q = inla.as.sparse(Q)
} else {
  ## faster. we only need to define the lower-triangular Q!
  i = c(
    ## diagonal
    1L, n, 2L:(n-1L),
    ## off-diagonal
    1L:(n-1L))
  j = c(
    ## diagonal
    1L, n, 2L:(n-1L),
    ## off-diagonal
    2L:n)
  x = param$prec/(1-param$rho^2) *
    c( ## diagonal
      1L, 1L, rep(1+param$rho^2, n-2L),
      ## off-diagonal
      rep(-param$rho, n-1L))
  Q = sparseMatrix(i=i, j=j, x=x, giveCsparse=FALSE)
}
return (Q)
}

```

Function mu

This function must return the mean of the model. Often, the mean is zero, but sometimes it might depend on the hyperparameters as well. If `numeric(0)` is returned, then this is equivalent that the mean is zero. An alternative in this example, would be to return `rep(0,n)`.

```

mu = function(n, theta)
{
  return(numeric(0))
}

```

Function log.norm.const

This function must return the log of the normalising constant. For the AR1-model the normalising constant is

$$\left(\frac{1}{\sqrt{2\pi}}\right)^n \tau_I^{n/2} (1 - \rho^2)^{1/2}$$

where

$$\tau_I = \tau / (1 - \rho^2).$$

The function can then be implemented as

```

log.norm.const = function(n, theta)
{
  ## return the log(normalising constant) for the model
  param = interpret.theta(n, theta)
  prec.innovation = param$prec / (1.0 - param$rho^2)
  val = n * (- 0.5 * log(2*pi) + 0.5 * log(prec.innovation)) + 0.5 * log(1.0 - param$rho^2)
  return (val)
}

```

NOTE: If the log-normalizing constant in any case need to be computed from scratch as

$$-\frac{n}{2} \log(2\pi) + \frac{1}{2} \log(|Q|),$$

then INLA will compute this if `numeric(0)` is returned, like

```
log.norm.const = function(n, theta)
{
  ## let INLA compute it
  return (numeric(0))
}
```

Function `log.prior`

This function must return the (log-)prior of the prior density for θ . For the AR1-model, we have for simplicity chosen this prior

$$\pi(\theta) = \Gamma(\exp(\theta_1); a, b) \exp(\theta_1) \times \mathcal{N}(\theta_2; \mu, \kappa)$$

so we can implement this as with our choices $a = b = 1$, $\mu = 0$ and $\kappa = 1$.

```
log.prior = function(n, theta)
{
  ## return the log-prior for the hyperparameters. the '+theta[1L]' is the log(Jacobian)
  ## for having a gamma prior on the precision and convert it into the prior for the
  ## log(precision).
  param = interpret.theta(n, theta)
  val = (dgamma(param$prec, shape = 1, rate = 1, log=TRUE) + theta[1L] +
        dnorm(theta[2L], mean = 0, sd = 1, log=TRUE))
  return (val)
}
```

Function `initial`

This function returns the initial values for θ .

```
initial = function(n, theta)
{
  ## return initial values
  ntheta = 2
  return (rep(1, ntheta))
}
```

Function `quit`

This function is called when all the computations are done and before exit-ing the C-program. Usually, there is nothing in particular to do, but if there is something that should be done, you can do this here.

```
quit = function(n, theta)
{
  return (invisible())
}
```

The complete definition of the AR1-model

For completeness, we include here the complete code for the AR1-model, collecting all the functions already defined. The function is predefined in the INLA-library.

```
'inla.rgeneric.ar1.model' = function(cmd = c("graph", "Q", "mu", "initial",
                                             "log.norm.const", "log.prior", "quit"),
                                     theta = NULL, args = NULL)
{
  ## this is an example of the 'rgeneric' model. here we implement
  ## the AR-1 model as described in inla.doc("ar1"), where 'rho' is
  ## the lag-1 correlation and 'prec' is the *marginal* (not
  ## conditional) precision.

  interpret.theta = function(n, theta)
  {
    ## internal helper-function to map the parameters from the internal-scale to the
    ## user-scale
    return (list(prec = exp(theta[1L]),
                 rho = 2*exp(theta[2L])/(1+exp(theta[2L])) - 1.0))
  }

  graph = function(n, theta)
  {
    ## return the graph of the model. the values of Q is only interpreted as zero or
    ## non-zero. return a sparse.matrix
    if (FALSE) {
      ## slow and easy: dense-matrices
      G = toeplitz(c(1, 1, rep(0, n-2L)))
      G = inla.as.sparse(G)
    } else {
      ## faster. we only need to define the lower-triangular of G
      i = c(
        ## diagonal
        1L, n, 2L:(n-1L),
        ## off-diagonal
        1L:(n-1L))
      j = c(
        ## diagonal
        1L, n, 2L:(n-1L),
        ## off-diagonal
        2L:n)
      x = 1 ## meaning that all are 1
      G = sparseMatrix(i=i, j=j, x=x, giveCsparse = FALSE)
    }
    return (G)
  }

  Q = function(n, theta)
  {
    ## returns the precision matrix for given parameters
    param = interpret.theta(n, theta)
    if (FALSE) {
      ## slow and easy: dense-matrices
      Q = param$prec/(1-param$rho^2) * toeplitz(c(1+param$rho^2, -param$rho, rep(0, n-2L)))
      Q[1, 1] = Q[n, n] = param$prec/(1-param$rho^2)
      Q = inla.as.sparse(Q)
    } else {
```

```

    ## faster. we only need to define the lower-triangular Q!
    i = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      1L:(n-1L))
    j = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      2L:n)
    x = param$prec/(1-param$rho^2) *
      c( ## diagonal
        1L, 1L, rep(1+param$rho^2, n-2L),
        ## off-diagonal
        rep(-param$rho, n-1L))
    Q = sparseMatrix(i=i, j=j, x=x, giveCsparse=FALSE)
  }
  return (Q)
}

mu = function(n, theta)
{
  return(numeric(0))
}

log.norm.const = function(n, theta)
{
  ## return the log(normalising constant) for the model
  param = interpret.theta(n, theta)
  prec.innovation = param$prec / (1.0 - param$rho^2)
  val = n * (- 0.5 * log(2*pi) + 0.5 * log(prec.innovation)) + 0.5 * log(1.0 - param$rho^2)
  return (val)
}

log.prior = function(n, theta)
{
  ## return the log-prior for the hyperparameters. the '+theta[1L]' is the log(Jacobian)
  ## for having a gamma prior on the precision and convert it into the prior for the
  ## log(precision).
  param = interpret.theta(n, theta)
  val = (dgamma(param$prec, shape = 1, rate = 1, log=TRUE) + theta[1L] +
    dnorm(theta[2L], mean = 0, sd = 1, log=TRUE))
  return (val)
}

initial = function(n, theta)
{
  ## return initial values
  ntheta = 2
  return (rep(1, ntheta))
}

quit = function(n, theta)
{
  return (invisible())
}

```

```
cmd = match.arg(cmd)
val = do.call(cmd, args = list(n = as.integer(args$n), theta = theta))
return (val)
}
```

Example of usage

```
n = 100
rho=0.9
x = arima.sim(n, model = list(ar = rho)) * sqrt(1-rho^2)
y = x + rnorm(n, sd = 0.1)
model = inla.rgeneric.define(inla.rgeneric.ar1.model, n=n)
formula = y ~ -1 + f(idx, model=model)
r = inla(formula, data = data.frame(y, idx = 1:n), family = "gaussian")
```


Example: the iid-model

The following function defines the iid-model, see `inla.doc("iid")`, which we give without further comments. To run this model in R, you may do `demo(rgeneric)`.

```
'inla.rgeneric.iid.model' = function(cmd = c("graph", "Q", "mu", "initial",
                                             "log.norm.const", "log.prior", "quit"),
                                     theta = NULL, args = NULL)
{
  ## this is an example of the 'rgeneric' model. here we implement the iid model as described
  ## in inla.doc("iid"), without the scaling-option

  interpret.theta = function(n, theta)
  {
    return (list(prec = exp(theta[1L])))
  }

  graph = function(n, theta)
  {
    G = Diagonal(n, x= rep(1, n))
    return (G)
  }

  Q = function(n, theta)
  {
    prec = interpret.theta(n, theta)$prec
    Q = Diagonal(n, x= rep(prec, n))
    return (Q)
  }

  mu = function(n, theta)
  {
    return(numeric(0))
  }

  log.norm.const = function(n, theta)
  {
    prec = interpret.theta(n, theta)$prec
    val = sum(dnorm(rep(0, n), sd = 1/sqrt(prec), log=TRUE))
    return (val)
  }

  log.prior = function(n, theta)
  {
    prec = interpret.theta(n, theta)$prec
    val = dgamma(prec, shape = 1, rate = 1, log=TRUE) + theta[1L]
    return (val)
  }

  initial = function(n, theta)
  {
    ntheta = 1
    return (rep(1, ntheta))
  }

  quit = function(n, theta)
  {
    return (invisible())
  }
}
```

```
}  
  
cmd = match.arg(cmd)  
val = do.call(cmd, args = list(n = as.integer(args$n), theta = theta))  
return (val)  
}
```

Example: a model for the mean structure

```
## In this example we do linear regression using 'rgeneric'.
## The regression model is  $y = a + b \cdot x + \text{noise}$ , and we
## define ' $a + b \cdot x + \text{tiny.noise}$ ' as a latent model.
## The dimension is length(x) and number of hyperparameters
## is 2 ('a' and 'b').
## This is a prototype example how similar situations
## could be approached, where essentially the latent model is a
## model for the 'mean' only.

rgeneric.linear.regression =
  function (cmd = c("graph", "Q", "mu", "initial", "log.norm.const",
                    "log.prior", "quit"),
           theta = NULL, args = NULL)
{
  ## artificial high precision to be added to the mean-model
  prec.high = exp(15)

  interpret.theta = function(theta) {
    return(list(a = theta[1L], b = theta[2L]))
  }

  graph = function(theta, x) {
    G = Diagonal(n = length(x), x=1)
    return(G)
  }

  Q = function(theta, x) {
    Q = prec.high * graph(theta, x)
    return(Q)
  }

  mu = function(theta, x) {
    par = interpret.theta(theta)
    return(par$a + par$b * x)
  }

  log.norm.const = function(theta, x) {
    ## the easiest is to let INLA compute this
    return(numeric(0))
  }

  log.prior = function(theta, x) {
    par = interpret.theta(theta)
    val = (dnorm(par$a, mean=0, sd = sqrt(1/0.001), log=TRUE) +
           dnorm(par$b, mean = 0, sd = sqrt(1/0.001), log=TRUE))
    return(val)
  }

  initial = function(theta, x) {
    return(rep(0, 2))
  }

  quit = function(theta, x) {
    return(invisible())
  }
}
```

```

    }

    cmd = match.arg(cmd)
    val = do.call(cmd, args = list(theta = theta, x = args$x))
    return(val)
}

a = 1
b = 2
n = 100
x = rnorm(n)
eta = a + b*x
y = eta + rnorm(n)

rgen = inla.rgeneric.define(model = rgeneric.linear.regression, x = x)
r = inla(y ~ -1 + f(idx, model=rgen),
        data = data.frame(y, idx = 1:n))
rr = inla(y ~ 1 + x,
        data = data.frame(y, x),
        control.fixed = list(prec.intercept = 0.001, prec = 0.001))

## compare the results with the 'truth'
par(mfrow=c(2, 1))
plot(r$marginals.hyperpar[['Theta1 for idx']], type="l", lwd=2, col="red",
     main = "Posterior for the intercept (red=rgeneric, blue=default)")
lines(rr$marginals.fixed$('Intercept'), lwd=2, col="blue")

plot(r$marginals.hyperpar[['Theta2 for idx']], type="l", lwd=2, col="red",
     main = "Posterior for the slope (red=rgeneric, blue=default)")
lines(rr$marginals.fixed$('x'), lwd=2, col="blue")

```