

Model `rgeneric`

This is a class of generic models allows the user to define latent model-component in `R`, for cases where the requested model is not yet implemented in `INLA`, and do the Bayesian inference using `INLA`. It will run slower as the model properties has to be evaluated in `R` within a `C`-program.

Defining a latent model in `R`

The use of this feature, is in short the following. First we pass our definition of the model `rmodel`, to define a `inla-rgeneric` object,

```
model = inla.rgeneric.define(rmodel, debug, ...)
```

Here, `rmodel` is model definition encoded as an `R`-function, `debug` is a logical parameter if debug information should be printed, and `...` are further named variables/environments that goes into the environment of `rmodel` (like dimension, prior-settings, etc). Then the model can be used as

```
y ~ ... + f(i, model=model)
```

Example: the AR1 model

The `rmodel` needs to follow some rules to provide the required features. As an example, we will show how to implement the AR1-model, see `inla.doc("ar1")`. This model is defined as¹

$$x_1 \sim \mathcal{N}(0, \tau)$$

and

$$x_t | x_1, \dots, x_{t-1} \sim \mathcal{N}(\rho x_{t-1}, \tau_I), \quad t = 2, \dots, n.$$

The scale-parameter is the marginal precision τ , but the joint density is more naturally expressed using the innovation precision $\tau_I = \tau/(1 - \rho^2)$. The joint density of x is Gaussian

$$\pi(x|\rho, \tau) = \left(\frac{1}{\sqrt{2\pi}}\right)^n \tau_I^{n/2} (1 - \rho^2)^{1/2} \exp\left(-\frac{\tau_I}{2} x^T R x\right)$$

where the precision-matrix is

$$Q = \tau_I R = \tau_I \begin{bmatrix} 1 & -\rho & & & & \\ -\rho & 1 + \rho^2 & -\rho & & & \\ & -\rho & 1 + \rho^2 & -\rho & & \\ & & \ddots & \ddots & \ddots & \\ & & & -\rho & 1 + \rho^2 & -\rho \\ & & & & -\rho & 1 \end{bmatrix}$$

There are two (hyper-)parameters for this model, it is the marginal precision τ and the lag-one correlation ρ . We reparameterise these as

$$\tau = \exp(\theta_1)$$

and

$$\rho = 2 \frac{\exp(\theta_2)}{1 + \exp(\theta_2)} - 1$$

It is required that the parameters $\theta = (\theta_1, \theta_2)$ have support on \mathfrak{R} and the priors for τ and ρ are given as the corresponding priors for θ_1 and θ_2 . **IMPORTANT: A good re-parameterisation is required**

¹The second argument in $\mathcal{N}(,)$ is the precision not the covariance.

for INLA to work well. A good parameterisation makes, ideally, the “Fisher information matrix” of θ constant (wrt to θ). It is sufficient to check this in a frequentistic setting with data directly from the AR(1) model, in this case. Note that INLA only provide the marginal posteriors for θ , but you can use `inla.tmarginal` to convert it to the appropriate marginals for ρ and τ .

We assign a (Gamma) $\Gamma(\cdot; a, b)$ prior (with mean a/b and variance a/b^2) for τ and a Gaussian prior $\mathcal{N}(\mu, \kappa)$ for θ_2 , so the joint prior for θ becomes

$$\pi(\theta) = \Gamma(\exp(\theta_1); a, b) \exp(\theta_1) \times \mathcal{N}(\theta_2; \mu, \kappa).$$

We will use $a = b = 1$, $\mu = 0$ and $\kappa = 1$.

In order to define the AR1-model, we need to make functions that returns

- the graph,
- the precision matrix $Q(\theta)$,
- the zero mean,
- the initial values of θ ,
- the log-normalising constant, and
- the log-prior

which except for the graph, depends on the current value of θ . We need to wrap this into a common function, which process the request from the C-program. The list of commands and its names

```
cmd = c("Q", "graph", "mu", "initial", "log.norm.const",
        "log.prior", "quit"),
```

are fixed. The skeleton-function for defining a model is

```
'inla.rgeneric.ar1.model' = function(
  cmd = c("graph", "Q", "mu", "initial", "log.norm.const",
          "log.prior", "quit"),
  theta = NULL)
{
  graph = function(theta){ <to be completed> }
  Q = function(theta) { <to be completed> }
  mu = function(theta) { <to be completed> }
  log.norm.const = function(theta) { <to be completed> }
  log.prior = function(theta) { <to be completed> }
  initial = function(theta) { <to be completed> }
  quit = function(theta) { <to be completed> }

  cmd = match.arg(cmd)
  val = do.call(cmd, args = list(theta = theta))
  return (val)
}
```

The input parameters are

cmd What to return

theta The values of the θ -parameters

Other parameters in the model definition, like n and possibly the parameters of the prior, goes into the “...” part of `inla.rgeneric.define`, like

```
model = inla.rgeneric.define(inla.rgeneric.ar1.model, n = 100)
```

and is assigned in the environment to `inla.rgeneric.ar1.model`. This is because several instances of `rgeneric` models will share the same `.GlobalEnv`.

Our next task, is the “fill in the blanks” in this function.

Function graph

This function must return a sparseMatrix, with the non-zero elements of the precision matrix. Only the lower-triangular part of the matrix is used.

```
graph = function(theta)
{
  ## return the graph of the model. the values of Q is only interpreted as zero or
  ## non-zero. return a sparse.matrix
  if (FALSE) {
    ## slow and easy: dense-matrices
    G = toeplitz(c(1, 1, rep(0, n-2L)))
    G = inla.as.sparse(G)
  } else {
    ## faster. we only need to define the lower-triangular of G
    i = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      1L:(n-1L))
    j = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      2L:n)
    x = 1 ## meaning that all are 1
    G = sparseMatrix(i=i, j=j, x=x, giveCsparse = FALSE)
  }
  return (G)
}
```

Function Q

This function must return the precision matrix $Q(\theta)$, and must be a sparseMatrix. Only the lower-triangular part of the matrix is used. We will make use of the helper function

```
interpret.theta = function(theta)
{
  ## internal helper-function to map the parameters from the internal-scale to the
  ## user-scale
  return (list(prec = exp(theta[1L]),
    rho = 2*exp(theta[2L])/(1+exp(theta[2L])) - 1.0))
}
```

to convert from θ_1 to τ , and from θ_2 to ρ . The Q-function can then be implemented as follows.

```
Q = function(theta)
{
  ## returns the precision matrix for given parameters
  param = interpret.theta(theta)
  if (FALSE) {
    ## slow and easy: dense-matrices
    Q = param$prec/(1-param$rho^2) * toeplitz(c(1+param$rho^2, -param$rho, rep(0, n-2L)))
    Q[1, 1] = Q[n, n] = param$prec/(1-param$rho^2)
    Q = inla.as.sparse(Q)
  } else {
    ## faster. we only need to define the lower-triangular Q!
    i = c(
      ## diagonal
```

```

      1L, n, 2L:(n-1L),
      ## off-diagonal
      1L:(n-1L))
  j = c(
    ## diagonal
    1L, n, 2L:(n-1L),
    ## off-diagonal
    2L:n)
  x = param$prec/(1-param$rho^2) *
    c( ## diagonal
      1L, 1L, rep(1+param$rho^2, n-2L),
      ## off-diagonal
      rep(-param$rho, n-1L))
  Q = sparseMatrix(i=i, j=j, x=x, giveCsparse=FALSE)
}
return (Q)
}

```

Function mu

This function must return the mean of the model. Often, the mean is zero, but sometimes it might depend on the hyperparameters as well. If `numeric(0)` is returned, then this is equivalent that the mean is zero. An alternative in this example, would be to return `rep(0,n)`.

```

mu = function(theta)
{
  return(numeric(0))
}

```

Function log.norm.const

This function must return the log of the normalising constant. For the AR1-model the normalising constant is

$$\left(\frac{1}{\sqrt{2\pi}}\right)^n \tau_I^{n/2} (1 - \rho^2)^{1/2}$$

where

$$\tau_I = \tau / (1 - \rho^2).$$

The function can then be implemented as

```

log.norm.const = function(theta)
{
  ## return the log(normalising constant) for the model
  param = interpret.theta(theta)
  prec.innovation = param$prec / (1.0 - param$rho^2)
  val = n * (- 0.5 * log(2*pi) + 0.5 * log(prec.innovation)) + 0.5 * log(1.0 - param$rho^2)
  return (val)
}

```

NOTE: If the log-normalizing constant in any case need to be computed from scratch as

$$-\frac{n}{2} \log(2\pi) + \frac{1}{2} \log(|Q|),$$

then INLA will compute this if `numeric(0)` is returned, like

```

log.norm.const = function(theta)
{
  ## let INLA compute it
  return (numeric(0))
}

```

Function log.prior

This function must return the (log-)prior of the prior density for θ . For the AR1-model, we have for simplicity chosen this prior

$$\pi(\theta) = \Gamma(\exp(\theta_1); a, b) \exp(\theta_1) \times \mathcal{N}(\theta_2; \mu, \kappa)$$

so we can implement this as with our choices $a = b = 1$, $\mu = 0$ and $\kappa = 1$.

```
log.prior = function(theta)
{
  ## return the log-prior for the hyperparameters. the '+theta[1L]' is the log(Jacobian)
  ## for having a gamma prior on the precision and convert it into the prior for the
  ## log(precision).
  param = interpret.theta(theta)
  val = (dgamma(param$prec, shape = 1, rate = 1, log=TRUE) + theta[1L] +
         dnorm(theta[2L], mean = 0, sd = 1, log=TRUE))
  return (val)
}
```

Function initial

This function returns the initial values for θ .

```
initial = function(theta)
{
  ## return initial values
  ntheta = 2
  return (rep(1, ntheta))
}
```

Function quit

This function is called when all the computations are done and before exit-ing the C-program. Usually, there is nothing in particular to do, but if there is something that should be done, you can do this here.

```
quit = function(theta)
{
  return (invisible())
}
```

The complete definition of the AR1-model

For completeness, we include here the complete code for the AR1-model, collecting all the functions already defined. The function is predefined in the INLA-library.

```
'inla.rgeneric.ar1.model' = function(cmd = c("graph", "Q", "mu", "initial",
                                             "log.norm.const", "log.prior", "quit"),
                                     theta = NULL)
{
  ## this is an example of the 'rgeneric' model. here we implement
  ## the AR-1 model as described in inla.doc("ar1"), where 'rho' is
  ## the lag-1 correlation and 'prec' is the *marginal* (not
  ## conditional) precision.

  interpret.theta = function(theta)
  {
    ## internal helper-function to map the parameters from the internal-scale to the
    ## user-scale
    return (list(prec = exp(theta[1L]),
```

```

        rho = 2*exp(theta[2L])/(1+exp(theta[2L])) - 1.0))
}

graph = function(theta)
{
  ## return the graph of the model. the values of Q is only interpreted as zero or
  ## non-zero. return a sparse.matrix
  if (FALSE) {
    ## slow and easy: dense-matrices
    G = toeplitz(c(1, 1, rep(0, n-2L)))
    G = inla.as.sparse(G)
  } else {
    ## faster. we only need to define the lower-triangular of G
    i = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      1L:(n-1L))
    j = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      2L:n)
    x = 1 ## meaning that all are 1
    G = sparseMatrix(i=i, j=j, x=x, giveCsparse = FALSE)
  }
  return (G)
}

Q = function(theta)
{
  ## returns the precision matrix for given parameters
  param = interpret.theta(theta)
  if (FALSE) {
    ## slow and easy: dense-matrices
    Q = param$prec/(1-param$rho^2) * toeplitz(c(1+param$rho^2, -param$rho, rep(0, n-2L)))
    Q[1, 1] = Q[n, n] = param$prec/(1-param$rho^2)
    Q = inla.as.sparse(Q)
  } else {
    ## faster. we only need to define the lower-triangular Q!
    i = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      1L:(n-1L))
    j = c(
      ## diagonal
      1L, n, 2L:(n-1L),
      ## off-diagonal
      2L:n)
    x = param$prec/(1-param$rho^2) *
      c( ## diagonal
        1L, 1L, rep(1+param$rho^2, n-2L),
        ## off-diagonal
        rep(-param$rho, n-1L))
    Q = sparseMatrix(i=i, j=j, x=x, giveCsparse=FALSE)
  }
}

```

```

    return (Q)
}

mu = function(theta)
{
  return(numeric(0))
}

log.norm.const = function(theta)
{
  ## return the log(normalising constant) for the model
  param = interpret.theta(theta)
  prec.innovation = param$prec / (1.0 - param$rho^2)
  val = n * (- 0.5 * log(2*pi) + 0.5 * log(prec.innovation)) + 0.5 * log(1.0 - param$rho^2)
  return (val)
}

log.prior = function(theta)
{
  ## return the log-prior for the hyperparameters. the '+theta[1L]' is the log(Jacobian)
  ## for having a gamma prior on the precision and convert it into the prior for the
  ## log(precision).
  param = interpret.theta(theta)
  val = (dgamma(param$prec, shape = 1, rate = 1, log=TRUE) + theta[1L] +
        dnorm(theta[2L], mean = 0, sd = 1, log=TRUE))
  return (val)
}

initial = function(theta)
{
  ## return initial values
  ntheta = 2
  return (rep(1, ntheta))
}

quit = function(theta)
{
  return (invisible())
}

cmd = match.arg(cmd)
val = do.call(cmd, args = list(theta = theta))
return (val)
}

```

Example of usage

```

n = 100
rho=0.9
x = arima.sim(n, model = list(ar = rho)) * sqrt(1-rho^2)
y = x + rnorm(n, sd = 0.1)
model = inla.rgeneric.define(inla.rgeneric.ar1.model, n=n)
formula = y ~ -1 + f(idx, model=model)
r = inla(formula, data = data.frame(y, idx = 1:n), family = "gaussian")

```

Example: the iid-model

The following function defines the iid-model, see `inla.doc("iid")`, which we give without further comments. To run this model in R, you may do `demo(rgeneric)`.

```
'inla.rgeneric.iid.model' = function(cmd = c("graph", "Q", "mu", "initial",
                                             "log.norm.const", "log.prior", "quit"),
                                     theta = NULL)
{
  ## this is an example of the 'rgeneric' model. here we implement the iid model as described
  ## in inla.doc("iid"), without the scaling-option

  interpret.theta = function(theta)
  {
    return (list(prec = exp(theta[1L])))
  }

  graph = function(theta)
  {
    G = Diagonal(n, x= rep(1, n))
    return (G)
  }

  Q = function(theta)
  {
    prec = interpret.theta(theta)$prec
    Q = Diagonal(n, x= rep(prec, n))
    return (Q)
  }

  mu = function(theta)
  {
    return(numeric(0))
  }

  log.norm.const = function(theta)
  {
    prec = interpret.theta(theta)$prec
    val = sum(dnorm(rep(0, n), sd = 1/sqrt(prec), log=TRUE))
    return (val)
  }

  log.prior = function(theta)
  {
    prec = interpret.theta(theta)$prec
    val = dgamma(prec, shape = 1, rate = 1, log=TRUE) + theta[1L]
    return (val)
  }

  initial = function(theta)
  {
    ntheta = 1
    return (rep(1, ntheta))
  }

  quit = function(theta)
  {
    return (invisible())
  }
}
```



```
}  
  
cmd = match.arg(cmd)  
val = do.call(cmd, args = list(theta = theta))  
return (val)  
}
```

Example: a model for the mean structure

```
## In this example we do linear regression using 'rgeneric'.
## The regression model is  $y = a + b \cdot x + \text{noise}$ , and we
## define ' $a + b \cdot x + \text{tiny.noise}$ ' as a latent model.
## The dimension is  $\text{length}(x)$  and number of hyperparameters
## is 2 ('a' and 'b').
## This is a prototype example how similar situations
## could be approached, where essentially the latent model is a
## model for the 'mean' only.

rgeneric.linear.regression =
  function(cmd = c("graph", "Q", "mu", "initial", "log.norm.const",
                  "log.prior", "quit"),
          theta = NULL)
{
  ## artificial high precision to be added to the mean-model
  prec.high = exp(15)

  interpret.theta = function(theta) {
    return(list(a = theta[1L], b = theta[2L]))
  }

  graph = function(theta) {
    G = Diagonal(n = length(x), x=1)
    return(G)
  }

  Q = function(theta) {
    Q = prec.high * graph(theta, x)
    return(Q)
  }

  mu = function(theta) {
    par = interpret.theta(theta)
    return(par$a + par$b * x)
  }

  log.norm.const = function(theta) {
    ## the easiest is to let INLA compute this
    return(numeric(0))
  }

  log.prior = function(theta) {
    par = interpret.theta(theta)
    val = (dnorm(par$a, mean=0, sd = sqrt(1/0.001), log=TRUE) +
           dnorm(par$b, mean = 0, sd = sqrt(1/0.001), log=TRUE))
    return(val)
  }

  initial = function(theta) {
    return(rep(0, 2))
  }

  quit = function(theta) {
    return(invisible())
  }
}
```

```

    cmd = match.arg(cmd)
    val = do.call(cmd, args = list(theta = theta))
    return(val)
}

a = 1
b = 2
n = 100
x = rnorm(n)
eta = a + b*x
y = eta + rnorm(n)

rgen = inla.rgeneric.define(model = rgeneric.linear.regression, x = x)
r = inla(y ~ -1 + f(idx, model=rgen),
        data = data.frame(y, idx = 1:n))
rr = inla(y ~ 1 + x,
        data = data.frame(y, x),
        control.fixed = list(prec.intercept = 0.001, prec = 0.001))

## compare the results with the 'truth'
par(mfrow=c(2, 1))
plot(r$marginals.hyperpar[['Theta1 for idx']], type="l", lwd=2, col="red",
     main = "Posterior for the intercept (red=rgeneric, blue=default)")
lines(rr$marginals.fixed$('Intercept'), lwd=2, col="blue")

plot(r$marginals.hyperpar[['Theta2 for idx']], type="l", lwd=2, col="red",
     main = "Posterior for the slope (red=rgeneric, blue=default)")
lines(rr$marginals.fixed$('x'), lwd=2, col="blue")

```